

AML_HW1_House Price Prediction

October 2, 2020

```
[ ]: # Team Members:  
#     Scarlett Huang (sh2557)  
#     Zihan Zhang (zz698)
```

```
[62]: # import modules  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
%matplotlib inline  
import seaborn as sbn  
from scipy.stats import norm  
from sklearn.preprocessing import StandardScaler  
import os
```

```
[64]: # import data  
train = pd.read_csv("train.csv")  
test = pd.read_csv("test.csv")  
  
# drop 'Id'  
Id = test['Id']  
train.drop('Id',axis=1,inplace=True)  
test.drop('Id',axis=1,inplace=True)
```

```
[35]: # explore data  
print (train.describe())  
print (train.columns)  
print (train.head(5))  
print (train.shape)  
print (test.shape)
```

```
<bound method NDFrame.describe of  
LotArea Street Alley LotShape \
```

	LotArea	Street	Alley	LotShape	MSSubClass	MSZoning	LotFrontage	
0	60	RL		65.0	8450	Pave	NaN	Reg
1	20	RL		80.0	9600	Pave	NaN	Reg
2	60	RL		68.0	11250	Pave	NaN	IR1
3	70	RL		60.0	9550	Pave	NaN	IR1
4	60	RL		84.0	14260	Pave	NaN	IR1
...

1455	60	RL	62.0	7917	Pave	NaN	Reg
1456	20	RL	85.0	13175	Pave	NaN	Reg
1457	70	RL	66.0	9042	Pave	NaN	Reg
1458	20	RL	68.0	9717	Pave	NaN	Reg
1459	20	RL	75.0	9937	Pave	NaN	Reg

	LandContour	Utilities	LotConfig	...	PoolArea	PoolQC	Fence	MiscFeature	\
0	Lvl	AllPub	Inside	...	0	NaN	NaN	NaN	
1	Lvl	AllPub	FR2	...	0	NaN	NaN	NaN	
2	Lvl	AllPub	Inside	...	0	NaN	NaN	NaN	
3	Lvl	AllPub	Corner	...	0	NaN	NaN	NaN	
4	Lvl	AllPub	FR2	...	0	NaN	NaN	NaN	
...	
1455	Lvl	AllPub	Inside	...	0	NaN	NaN	NaN	
1456	Lvl	AllPub	Inside	...	0	NaN	MnPrv	NaN	
1457	Lvl	AllPub	Inside	...	0	NaN	GdPrv	Shed	
1458	Lvl	AllPub	Inside	...	0	NaN	NaN	NaN	
1459	Lvl	AllPub	Inside	...	0	NaN	NaN	NaN	

	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	0	2	2008	WD	Normal	208500
1	0	5	2007	WD	Normal	181500
2	0	9	2008	WD	Normal	223500
3	0	2	2006	WD	Abnorml	140000
4	0	12	2008	WD	Normal	250000
...
1455	0	8	2007	WD	Normal	175000
1456	0	2	2010	WD	Normal	210000
1457	2500	5	2010	WD	Normal	266500
1458	0	4	2010	WD	Normal	142125
1459	0	6	2008	WD	Normal	147500

[1460 rows x 80 columns]>

```
Index(['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley',
      'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
      'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
      'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
      'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
      'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
      'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
      'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
      'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
      'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
      'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
      'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
      'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
      'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
      'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
```

```

        'MoSold', 'YrSold', 'SaleType', 'SaleCondition', 'SalePrice'],
        dtype='object')
MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0          60      RL          65.0      8450   Pave   NaN      Reg
1          20      RL          80.0      9600   Pave   NaN      Reg
2          60      RL          68.0     11250   Pave   NaN      IR1
3          70      RL          60.0      9550   Pave   NaN      IR1
4          60      RL          84.0     14260   Pave   NaN      IR1

LandContour Utilities LotConfig ... PoolArea PoolQC Fence MiscFeature \
0          Lvl     AllPub    Inside ...        0    NaN    NaN        NaN
1          Lvl     AllPub    FR2     ...        0    NaN    NaN        NaN
2          Lvl     AllPub    Inside ...        0    NaN    NaN        NaN
3          Lvl     AllPub    Corner ...        0    NaN    NaN        NaN
4          Lvl     AllPub    FR2     ...        0    NaN    NaN        NaN

MiscVal MoSold  YrSold  SaleType  SaleCondition  SalePrice
0         0        2   2008        WD          Normal    208500
1         0        5   2007        WD          Normal    181500
2         0        9   2008        WD          Normal    223500
3         0        2   2006        WD      Abnorml     140000
4         0       12   2008        WD          Normal    250000

```

```

[5 rows x 80 columns]
(1460, 80)
(1459, 80)

```

```

[36]: # analyze target
      target = 'SalePrice'
      train[target].describe()

```

```

[36]: count      1460.000000
      mean      180921.195890
      std       79442.502883
      min       34900.000000
      25%      129975.000000
      50%      163000.000000
      75%      214000.000000
      max       755000.000000
      Name: SalePrice, dtype: float64

```

```


[37]: # use seaborn to plot SalePrice to see its distribution
      sbn.distplot(train[target])
      # result resembles normal distribution

```

```

[37]: <matplotlib.axes._subplots.AxesSubplot at 0x7f94964f7fa0>

```




```
[38]: # calculate kurtosis and skewness of the plot above
print ("skewness=",train[target].skew())
print ("kurtosis=",train[target].kurt())
```

```
skewness= 1.8828757597682129
kurtosis= 6.536281860064529
```

```
[39]: # perform the logarithm operation to make the data distribution of target
      ↪function more in line with the standard normal distribution
train[target] = np.log1p(train[target])
sbn.distplot(train[target])
```

```
[39]: <matplotlib.axes._subplots.AxesSubplot at 0x7f94964f0df0>
```




```
[40]: # Data merging
all_data= pd.concat([train,test],axis=0,join='outer',ignore_index=True)
all_data.drop([target], axis=1, inplace=True)
all_data.shape
```

```
[40]: (2919, 79)
```

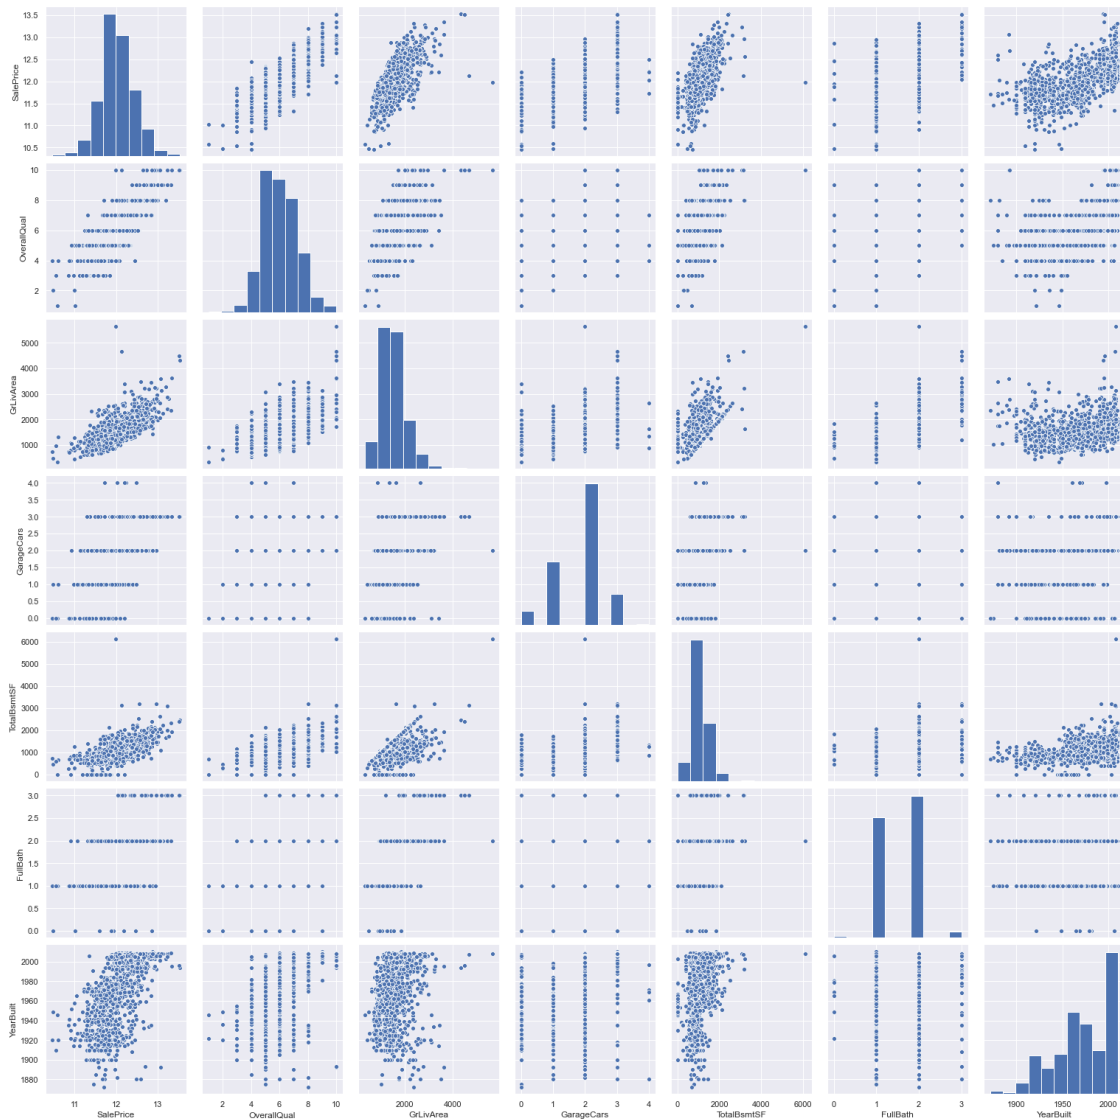
```
[41]: # use sns.heatmap() to conduct correlation analysis between features and target
plt.subplots(figsize=(24,24))
sns.heatmap(train.corr(),square=True,linecolor='white',linewidths=1, annot=True)
```

```
[41]: <matplotlib.axes._subplots.AxesSubplot at 0x7f93cc665bb0>
```



```
[42]: # identified Top10 most related features from the heatmap above: 'OverallQual',
      ↪ 'GrLivArea', 'GarageCars', 'GarageArea', 'TotalBsmntSF', '1stFlrSF',
      ↪ 'FullBath', 'YearBuilt', 'YearRemodAdd'
      # these pairs of features can be merged into one:
      ↪ 'GarageCars'+ 'GarageArea' 'TotalBsmntSF'+ '1stFlrSF' 'YearBuilt'+ 'YearRemodAdd'
```

```
[43]: # observe the correlation tendency between each feature
sbn.set()
cols = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmntSF',
      ↪ 'FullBath', 'YearBuilt']
sbn.pairplot(train[cols], height = 3)
plt.show()
# result shows that the six features are all positively correlated with target
```



```
[44]: # data processing & cleaning
all_data.shape
print('MSSubClass.type:', all_data.MSSubClass.dtypes)
all_data.MSSubClass=all_data.MSSubClass.astype(str)
all_data.MSSubClass.value_counts()
```

MSSubClass.type: int64

```
[44]: 20      1079
      60      575
      50      287
      120     182
      30      139
```

```

70      128
160     128
80      118
90      109
190      61
85      48
75      23
45      18
180     17
40       6
150      1
Name: MSSubClass, dtype: int64

```

```

[45]: # use pandas get_dummies to tranform categorical features into numerical
      ↪ features
all_data = pd.get_dummies(all_data)

# fill the missing values with the mean from each column
cols_mean=all_data.mean()
all_data=all_data.fillna(cols_mean)
all_data.isnull().sum().sum()
all_data.head()

```

```

[45]:   LotFrontage  LotArea  OverallQual  OverallCond  YearBuilt  YearRemodAdd  \
0          65.0    8450           7           5        2003        2003
1          80.0    9600           6           8        1976        1976
2          68.0   11250           7           5        2001        2002
3          60.0    9550           7           5        1915        1970
4          84.0   14260           8           5        2000        2000

      MasVnrArea  BsmtFinSF1  BsmtFinSF2  BsmtUnfSF  ...  SaleType_ConLw  \
0          196.0    706.0        0.0    150.0  ...              0
1           0.0    978.0        0.0    284.0  ...              0
2          162.0    486.0        0.0    434.0  ...              0
3           0.0    216.0        0.0    540.0  ...              0
4          350.0    655.0        0.0    490.0  ...              0

      SaleType_New  SaleType_Oth  SaleType_WD  SaleCondition_Abnorml  \
0              0              0              1              0
1              0              0              1              0
2              0              0              1              0
3              0              0              1              1
4              0              0              1              0

      SaleCondition_AdjLand  SaleCondition_Alloca  SaleCondition_Family  \
0              0              0              0
1              0              0              0

```


2	0	0	0
3	0	0	0
4	0	0	0

	SaleCondition_Normal	SaleCondition_Partial
0	1	0
1	1	0
2	1	0
3	0	0
4	1	0

[5 rows x 303 columns]

```
[46]: # standardize numeric variables
numer_cols=all_data.columns[all_data.dtypes != 'object']
numer_cols_mean=all_data.loc[:,numer_cols].mean()
numer_cols_std=all_data.loc[:,numer_cols].std()
all_data.loc[:,numer_cols]=(all_data.loc[:,numer_cols]-numer_cols_mean)/
    ↪numer_cols_std
```

```
[47]: # separate training set and testing set
dummy_train = all_data.loc[train.index]
dummy_test = all_data.loc[test.index]
dummy_train.shape
```

[47]: (1460, 303)

```
[48]: # do the modeling
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.kernel_ridge import KernelRidge
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
```

```
[49]: x_train = all_data[:train.shape[0]]
x_test = all_data[train.shape[0]:]
y_train = train.SalePrice
```

```
[50]: # Use linear regression model based on least squares, and turn it into a
    ↪quadratic polynomial model through PolynomialFeatures

lr_model = Pipeline([
    ('poly', PolynomialFeatures(degree=2)),
    ('lr', LinearRegression())
])
```

```

lr_model.fit(x_train, y_train)
lr_predict = cross_val_predict(lr_model, x_train, y_train, verbose=True,
    ↪n_jobs=-1, cv=3)
lr_mse = mean_squared_error(lr_predict, y_train)
lr_score = np.sqrt(lr_mse)
print("linear regression score: ", lr_score)

lr_model.fit(x_train, y_train)
lr_preds = np.expm1(lr_model.predict(x_test))
lr_solution = pd.DataFrame({"id":Id, target:lr_preds})
lr_solution.to_csv("lr_output.csv", index = False)

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 15.2s finished

linear regression score: 0.19982095204644656

[59]: *# Use Ridge Regression (regularized least squares-L2 regularization)*

```

ridge_model = KernelRidge(degree=2, alpha=0.08, kernel='polynomial')
ridge_predict = cross_val_predict(ridge_model, x_train, y_train, cv=3,
    ↪verbose=True, n_jobs=-1)
ridge_mse = mean_squared_error(y_train, ridge_predict)
ridge_score = np.sqrt(ridge_mse)
print("ridge regression score: ", ridge_score)

ridge_model.fit(x_train, y_train)
ridge_preds = np.expm1(ridge_model.predict(x_test))
ridge_solution = pd.DataFrame({"id":Id, target:ridge_preds})
ridge_solution.to_csv("ridge_output.csv", index = False)

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 0.1s finished

ridge regression score: 0.14268299026580086

[57]: *# Use Lasso Regression (regularized least squares-L1 regularization)*

```

lasso_model = Lasso(alpha=0.0008, random_state=1, max_iter=5000)
lasso_predict = cross_val_predict(lasso_model, x_train, y_train, cv=3,
    ↪verbose=True, n_jobs=-1)
lasso_mse = mean_squared_error(lasso_predict, y_train)
lasso_score = np.sqrt(lasso_mse)
print("lasso regression score: ", lasso_score)

lasso_model.fit(x_train, y_train)
lasso_preds = np.expm1(lasso_model.predict(x_test))
lasso_solution = pd.DataFrame({"id":Id, target:lasso_preds})

```

```
lasso_solution.to_csv("lasso_output.csv", index = False)
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 3 out of 3 | elapsed: 0.1s finished
```

```
lasso regression score: 0.14412691102769676
```

```
[53]: # Briefly describe your findings from training regularized and unregularized
      ↪ models.

      # 1.unregularized models - linear regression (least squares):
      # 1) exists overfitting problem.

      # 2.regularized models - ridge regression & lasso regression:
      # 1) when  $\alpha$  is setted to be small enough, it reduces risk of overfitting
      # and also maintains the complexity of the model.
      # 2) L1 regularization is useful in alleviating the overfitting problem,
      # but it may also cause the loss of precision and the problem of
      ↪ insufficient generalization ability.
      # 3) also exist a certain degree of overfitting.
      # 4) show better performance than the unregularized model.
      # 5) the scores of Lasso (L1) and Ridge (L2) appear to be very close to each
      ↪ other.
```

```
[54]: # The result is submitted to Kaggle.

      # Kaggle Link: https://www.kaggle.com/scarletty
```